

#Robot Learning Assignment 01

Team members:

Roberto Cai Wu

Ramesh Kumar

In [155]:

```
import numpy as np
import matplotlib.pyplot as plt
import random
from random import randint
```

Exercise 1.1

In [107]:

```
# To compute individual expected value of each arm
arms = np.array([[2,3],[-1,5],[1,5],[-2,4],[0,3],[2,6]])
expected_value_each_arm = np.mean(arms,axis=1)
for i in range(6):
    print('expected_value arm ' + str(i+1) + " = ",expected_value_each_arm[i])
cumulative_expected_value = np.sum(expected_value_each_arm)
print('cumulative expected value',cumulative_expected_value)
```

```
expected_value arm 1 = 2.5
expected_value arm 2 = 2.0
expected_value arm 3 = 3.0
expected_value arm 4 = 1.0
expected_value arm 5 = 1.5
expected_value arm 6 = 4.0
cumulative expected value 14.0
```

The expected reward for sampling the 6 arms is:

$$E[R] = \sum_{k=0}^{N=5} r_k$$

, where

$$r_k$$

is the expected value for choosing arm k .

In this case we have:

$$R = r_0 + r_1 + r_2 + r_3 + r_4 + r_5 = 14.0$$

In [150]:

```

class Armed_Bandits(object):
    def __init__(self, bandits):
        self.bandits = bandits #Receives the intervals
        self.arms = len(bandits[:,0]) #Number of arms

    # Exercise 1.2

    # Return a random uniform number from the selected
    def sample_bandit(self, bandit_index):
        expected_value = np.random.uniform(self.bandits
        return expected_value

    #Average expected reward from taking uniform sample
    def average_reward(self, iterations):
        self.Q = np.zeros(self.arms)
        self.C_actions = np.zeros(self.arms)
        for i in range(iterations):
            arm_index = np.random.randint(0, self.arms)
            self.C_actions[arm_index] += 1
            sample = self.sample_bandit(arm_index)
            self.Q[arm_index] += (sample - self.Q[arm_i
            best_arm = np.argmax(self.Q)
            best_reward = self.Q[best_arm]
            cumulative_reward = np.sum(self.Q)
            print("Arm with highest reward = ", best_arm, " w
            print("Cumulative rewards = ", cumulative_reward
            print(self.Q)
            print(self.C_actions)

    # Exercise 1.3

    def print_Q(self):
        percentage = 100*(self.C_actions/np.sum(self.C_

        print("\nRewards: ", self.Q)
        print("Number of actions: ", self.C_actions)
        print("Percentages: ", percentage)
        print(np.sum(self.C_actions))

    #Expected reward from taking an e-greedy rate
    def e_greedy(self, iterations, frequency, e):
        self.Q = np.zeros(self.arms)

```

```

self.C_actions = np.zeros(self.arms)
for i in range(iterations):
    # 90% exploitation rate
    if np.random.random() > e:
        arm_index = np.argmax(self.Q)
    # 10% exploration rate
    else:
        arm_index = np.random.randint(0, self.arms)
    self.C_actions[arm_index] += 1
    sample = self.sample_bandit(arm_index)
    self.Q[arm_index] += (sample - self.Q[arm_index])

    if ((i + 1) % frequency) == 0:
        self.print_Q()

```

Exercise 1.4

#Expected reward adding a learning rate

```

def third_arm_uniformly(self, iterations, frequency, e):
    self.Q = np.zeros(self.arms)
    self.C_actions = np.zeros(self.arms)
    for i in range(iterations):
        # 90% exploitation rate
        if np.random.random() > e:
            arm_index = np.argmax(self.Q)
        # 10% exploration rate
        else:
            arm_index = np.random.randint(0, self.arms)
        self.C_actions[arm_index] += 1
        sample = self.sample_bandit(arm_index)
        # Calculate Q with learning rate alpha
        self.Q[arm_index] += alpha * (sample - self.Q[arm_index])

        if ((i + 1) % frequency) == 0:
            self.print_Q()

    # after 500 iterations
    if i == 500:
        self.bandits[2] = np.array([6, 8])

```

Exercise 1.5

```

def print_optimistic_Q(self):
    percentage = 100*(self.C_actions/np.sum(self.C_
    print("\nRewards: ",self.optimistic_Q)
    print("Number of actions: ",self.C_actions)
    print("Percentages: ",percentage)
    print(np.sum(self.C_actions))

#Expected reward with  $Q_k = 10$ 
def optimistic_intiallization(self,iterations,frequ
#Starting  $\bar{Q}$  with 10
    self.optimistic_Q = np.ones(self.arms)*10
    self.C_actions = np.zeros(self.arms)
    for i in range(iterations):
        # 90% exploitation rate
        if np.random.random() > e:
            arm_index = np.argmax(self.optimistic_Q
        # 10% exploration rate
        else:
            arm_index = np.random.randint(0,self.ar
            self.C_actions[arm_index] += 1
            sample = self.sample_bandit(arm_index)
            # Calculate Q with learning rate alpha
            self.optimistic_Q[arm_index] += alpha * (sa

        if ((i + 1) % frequency) == 0:
            self.print_optimistic_Q()

    # after 500 iterations
    if i == 500:
        self.bandits[2] = np.array([6,8])

```

Exercise 1.2

In [151]:

```
reward_distribution = np.array([[2,3],
                                [-1,5],
                                [1,5],
                                [-2,4],
                                [0,3],
                                [2,6]])
bandits = Armed_Bandits(reward_distribution)
print('\n With 100 iterations:')
samples = 100
bandits.average_reward(samples)
print('\n With 1000 iterations:')
samples = 1000
bandits.average_reward(samples)
```

With 100 iterations:
Arm with highest reward = 5 with reward
= 3.64987684529
Cumulative rewards = 13.7677137739
[2.39178293 1.75197068 3.16878242 1.18
274535 1.62255555 3.64987685]
[19. 13. 15. 18. 12. 23.]

With 1000 iterations:
Arm with highest reward = 5 with reward
= 3.98849971899
Cumulative rewards = 13.8512897643
[2.47302757 2.06342472 2.89355753 0.96
597462 1.4668056 3.98849972]
[173. 183. 147. 170. 165. 162.]

We can see that our cumulative reward is different from our expected reward from Ex. 1.1 because 100 iterations are not enough to converge to expected reward. If we increase the number of iterations, it will come closer to the expected rewards.

Exercise 1.3

In [153]:

```
samples = 1000  
print_at = 100  
e = 0.1  
bandits.e_greedy(samples, print_at, e)
```

```
Rewards: [ 2.49308515  1.11773078  2.4708
2835 -0.37679751  1.64625208  3.96322189]
Number of actions: [ 31.   2.   3.   1.
 3.  60.]
Percentages: [ 31.   2.   3.   1.   3.  6
0.]
100.0
```

```
Rewards: [ 2.49308515  1.21364114  2.6464
9659 0.97825969  1.3974162   4.06109198]
Number of actions: [ 31.   3.   6.
 2.   7. 151.]
Percentages: [ 15.5   1.5   3.    1.
 3.5 75.5]
200.0
```

```
Rewards: [ 2.48977484  1.21364114  2.9929
3676 0.97825969  1.3974162   4.03000485]
Number of actions: [ 34.   3.   9.
 2.   7. 245.]
Percentages: [ 11.33333333  1.
 3.    0.66666667  2.33333333
 81.66666667]
300.0
```

```
Rewards: [ 2.49453413  1.62334728  2.9495
6422 1.1127994   1.3974162   4.00396897]
Number of actions: [ 36.   4.  10.
 5.   7. 338.]
Percentages: [ 9.    1.    2.5   1.25
 1.75 84.5 ]
400.0
```

```
Rewards: [ 2.49481667  0.85149908  2.9495
6422 1.18812529  1.35581282  3.98568296]
Number of actions: [ 38.   7.  10.
 6.  10. 429.]
Percentages: [ 7.6   1.4   2.    1.2
 2.   85.8]
500.0
```

```
Rewards: [ 2.49274703  1.4877718  3.2333
4982  1.62004036  1.36266219  3.9545445 ]
Number of actions: [ 41.  10.  14.
9.  11.  515.]
Percentages: [ 6.83333333  1.66666667
2.33333333  1.5  1.83333333
85.83333333]
600.0
```

```
Rewards: [ 2.48540142  1.89892538  3.1759
1304  1.94480829  1.20167126  3.92558516]
Number of actions: [ 44.  12.  16.  1
1.  14.  603.]
Percentages: [ 6.28571429  1.71428571
2.28571429  1.57142857  2.  86.
14285714]
700.0
```

```
Rewards: [ 2.49541177  1.81344415  3.1759
1304  1.87291406  1.34142261  3.91730631]
Number of actions: [ 47.  13.  16.  1
2.  18.  694.]
Percentages: [ 5.875  1.625  2.
1.5  2.25  86.75 ]
800.0
```

```
Rewards: [ 2.49227146  1.81344415  3.2813
86  1.60476038  1.40334237  3.90823125]
Number of actions: [ 50.  13.  19.  1
4.  19.  785.]
Percentages: [ 5.55555556  1.44444444
2.11111111  1.55555556  2.11111111
87.22222222]
900.0
```

```
Rewards: [ 2.50366081  1.89894946  3.2185
6378  1.60476038  1.44771943  3.92464467]
Number of actions: [ 52.  16.  24.  1
4.  20.  874.]
Percentages: [ 5.2  1.6  2.4  1.4
2.  87.4]
1000.0
```

From the exercise we can observe that taking an e-greedy approach quickly recognizes arm 6 as the optimal reward.

Exercise 1.4

In [154]:

```
samples = 1000  
print_at = 100  
e = 0.1  
alpha = 0.01  
bandits.third_arm_uniformly(samples, print_at, e, alpha)
```

```
Rewards: [ 1.54491342  0.06676239  0.
          0.03814428  0.05344189  0.07578076]
Number of actions: [ 90.   2.   0.   1.
                    5.   2.]
Percentages: [ 90.   2.   0.   1.   5.
              2.]
100.0
```

```
Rewards: [ 2.11924825  0.20518657  0.0848
9823  0.05439157  0.05434393  0.10604299]
Number of actions: [ 181.   6.   2.
                    2.   6.   3.]
Percentages: [ 90.5   3.   1.   1.
              3.   1.5]
200.0
```

```
Rewards: [ 2.34780504  0.22262436  0.0848
9823  0.05439157  0.08173925  0.12873107]
Number of actions: [ 276.   7.   2.
                    2.   9.   4.]
Percentages: [ 92.                2.33333333
              0.66666667  0.66666667  3.                1.
              33333333]
300.0
```

```
Rewards: [ 2.43247169  0.27905352  0.1090
9016  0.05439157  0.11454337  0.21514756]
Number of actions: [ 367.  10.   3.
                    2.  12.   6.]
Percentages: [ 91.75   2.5   0.75  0.5
              3.   1.5 ]
400.0
```

```
Rewards: [ 2.47664069  0.30835028  0.1381
9775  0.0284417  0.13252977  0.26957362]
Number of actions: [ 458.  12.   4.
                    4.  14.   8.]
Percentages: [ 91.6   2.4   0.8   0.8
              2.8   1.6]
500.0
```

```
Rewards: [ 2.48531833  0.30257881  0.4378
0551  0.0284417   0.13426864  0.32413685]
Number of actions: [ 551.   13.    8.
 4.   15.   9.]
Percentages: [ 91.83333333  2.16666667
1.33333333  0.66666667  2.5          1.
5          ]
600.0
```

```
Rewards: [ 2.48092838  0.29030607  0.5708
7771  0.06261835  0.14654593  0.41205994]
Number of actions: [ 642.   14.   10.
 6.   16.   12.]
Percentages: [ 91.71428571  2.
1.42857143  0.85714286  2.28571429
1.71428571]
700.0
```

```
Rewards: [ 2.4658027   0.34684495  0.5708
7771  0.06261835  0.14654593  0.46979247]
Number of actions: [ 738.   16.   10.
 6.   16.   14.]
Percentages: [ 92.25   2.      1.25   0.75
2.      1.75]
800.0
```

```
Rewards: [ 2.45433359  0.37328475  0.6444
1865  0.05832265  0.16402196  0.46979247]
Number of actions: [ 832.   18.   11.
 8.   17.   14.]
Percentages: [ 92.44444444  2.
1.22222222  0.88888889  1.88888889
1.55555556]
900.0
```

```
Rewards: [ 2.44684267  0.37328475  0.7072
7947  0.09710139  0.16402196  0.59530486]
Number of actions: [ 924.   18.   12.   1
1.   17.   18.]
Percentages: [ 92.4   1.8   1.2   1.1
1.7   1.8]
1000.0
```

From 1000 iterations is not possible to notice a change, but if we increase the learning rate or the number of iterations, we can see that the reward for the third arm will slowly increase. Since the learning rate is very slow, it initially fails to recognize arm 6 as the best arm.

Exercise 1.5

In [146]:

```
samples = 1000  
print_at = 100  
e = 0.1  
bandits.optimistic_intiallization(samples, print_at, e, al
```

```
Rewards: [ 9.1460399  9.05407557  9.1353
2991 9.1402683  9.14047616  9.10418347]
Number of actions: [ 12.  15.  35.  11.
 11. 16.]
Percentages: [ 12.  15.  35.  11.  11.  1
6.]
100.0
```

```
Rewards: [ 8.37082428  8.37018609  8.4013
3363 8.32877368  8.35074615  8.36017909]
Number of actions: [ 24.  24.  77.  21.
 22. 32.]
Percentages: [ 12.  12.  38.5  10.5  1
1. 16. ]
200.0
```

```
Rewards: [ 7.81718793  7.77105423  7.8641
895 7.83495586  7.81781261  7.86131371]
Number of actions: [ 34.  34.  132.  2
7. 30. 43.]
Percentages: [ 11.33333333  11.33333333
44. 9. 10. 1
4.33333333]
300.0
```

```
Rewards: [ 7.40490364  7.40336803  7.4564
4529 7.45648522  7.40331951  7.44878322]
Number of actions: [ 42.  40.  194.  3
2. 37. 55.]
Percentages: [ 10.5  10.  48.5  8.
9.25 13.75]
400.0
```

```
Rewards: [ 7.21670985  7.1532649  7.2426
5747 7.17958278  7.23068591  7.21510783]
Number of actions: [ 46.  45.  269.  3
7. 40. 63.]
Percentages: [ 9.2  9.  53.8  7.4
8. 12.6]
500.0
```

```
Rewards: [ 6.97478672  6.98274777  7.0624
5285  6.94312021  7.03490721  7.0222882 ]
Number of actions: [  51.   49.  348.   4
1.   43.  68.]
Percentages: [  8.5           8.16666667
58.           6.83333333  7.16666667
11.33333333]
600.0
```

```
Rewards: [ 6.83933706  6.87995323  6.9944
6477  6.85317733  6.84075258  6.97283102]
Number of actions: [  54.   51.  435.   4
3.   47.  70.]
Percentages: [  7.71428571  7.28571429
62.14285714  6.14285714  6.71428571  1
0.           ]
700.0
```

```
Rewards: [ 6.7158872  6.67937351  6.9743
4241  6.61884168  6.74350814  6.88559819]
Number of actions: [  57.   55.  520.   4
6.   49.  73.]
Percentages: [  7.125   6.875  65.
5.75   6.125  9.125]
800.0
```

```
Rewards: [ 6.62978728  6.6173423  7.0038
4964  6.47342693  6.64626018  6.88559819]
Number of actions: [  59.   56.  613.   4
8.   51.  73.]
Percentages: [  6.55555556  6.22222222
68.11111111  5.33333333  5.66666667
8.11111111]
900.0
```

```
Rewards: [ 6.54503904  6.57929354  6.9780
5827  6.39914874  6.48062454  6.84033098]
Number of actions: [  61.   57.  705.   4
9.   54.  74.]
Percentages: [  6.1   5.7  70.5   4.9
5.4   7.4]
1000.0
```

Since, initial action values can be used as a simple way of encouraging exploration. Therefore, we have assigned initial values to 10.0.

Furthermore, We can see that the selected arm is constantly changing because the higher initial reward is lowering after every selection. But overall, it will still converge slowly to the optimal arm.

In []:

In []: