

Robot Learning

Assignment 5

Due Tuesday, June 13th, before class.

5.1)

Develop an implementation of the game BeCareful, a simplified version of Black-jack:

- The game is played with an infinite deck of cards (i.e. cards are sampled with replacement)
- Each draw from the deck results in a value between 1 and 10 (uniformly distributed) with a color of red (probability 1/3) or black (probability 2/3).
- There are no aces or picture (face) cards in this game
- At the start of the game both the player and the dealer draw one black card (fully observed)
- Each turn the player may either stick or hit
- If the player hits then he or she draws another card from the deck
- If the player sticks he or she receives no further cards
- The values of the player's cards are added (black cards) or subtracted (red cards)
- If the player's sum exceeds 21, or becomes less than 1, then she "goes bust" and loses the game (reward -1)
- If the player sticks then the dealer starts taking turns. The dealer always sticks on any sum of 17 or greater, and hits otherwise.
- If the dealer goes bust, then the player wins; otherwise, the outcome { win (reward +1), lose (reward -1), or draw (reward 0) } is the player with the largest sum.

Specifically, write a function $(s', r) = \text{advance}(s, a)$, which takes as input a state s (dealer's first card 1–10 and the player's sum 1–21), and an action a (hit or stick), and returns a sample of the next state s' (which may be terminal if the game is finished) and reward $r \in \{1, 0, -1\}$ for winning, draw, and losing. All non-terminal rewards are zero. There is no discounting ($\gamma = 1$). You should treat the dealer's moves as part of the environment, i.e. calling $\text{advance}(s, \text{stick})$ will play out the dealer's cards and return the final reward and terminal state.

6 points

5.2)

Implement Sarsa(λ) for BeCareful. Initialise the value function $Q(s, a)$ to zero. Use a time-varying scalar step-size of $\alpha_t = 1/N(s_t, a_t)$ and an ϵ -greedy exploration strategy with $\epsilon_t = N_0/(N_0 + N(s_t))$, where $N_0 = 50$ is a constant, $N(s)$ is the number of times state s has been visited, and $N(s, a)$ is the number of times action a has been selected from state s .

Run the algorithm with parameter values $\lambda \in \{0, 0.1, 0.2, \dots, 1\}$. Stop exploration and learning after 1000 episodes and plot the accumulated reward for the next 100 episodes against λ .

7 points

5.3)

Use now a simple coarse coding value function approximator that is based on a binary feature vector $\phi(s, a)$ with $3 * 6 * 2 = 36$ features. Each binary feature has a value of 1 iff (s, a) lies within the cuboid of state-space corresponding to that feature, and the action corresponding to that feature. The cuboids have the following overlapping intervals:

dealer(s) = {[1, 4], [4, 7], [7, 10]}

player(s) = {[1, 6], [4, 9], [7, 12], [10, 15], [13, 18], [16, 21]}

a = {hit, stick}

where

- dealer(s) is the value of the dealer's first card (1–10)
- sum(s) is the sum of the player's cards (1–21)

Repeat the Sarsa(λ) experiment from 5.2), but using linear value function approximation $Q(s, a) = \phi(s, a)^T \theta$. Use a constant exploration of $\epsilon = 0.05$ and a constant step-size of $\alpha = 0.01$.

Again, for $\lambda \in \{0, 0.1, 0.2, \dots, 1\}$ learn from 1000 episodes, stop learning and exploration and plot the accumulated reward for the next 100 episodes against λ .

7 points